U.S. BANK

Risk Management & Fraud Operations Policy

Sample Policy Document for AI/RAG Hackathon Demonstration

Document ID	RISK-FRAUD-POLICY-SAMPLE-2026
Version	1.0 (Mock)
Effective Date	2026-01-01
Policy Owner	Fraud Operations & Risk Management
Classification	Internal - Demonstration Only

IMPORTANT: This is a synthetic sample policy created for demonstration and RAG testing. It is not an official U.S. Bank policy and must not be used for real-world fraud, compliance, credit, or operational decisions.

Purpose of this sample

This document intentionally contains clearly defined policy sections, thresholds, conditions, exceptions, required actions, and adjacent contextual material so that a retrieval-augmented generation (RAG) system can retrieve precise evidence and cite the appropriate policy section.

Table of Contents

1. Purpose and Scope

This policy establishes minimum fraud-risk review requirements for suspicious consumer credit-card transactions. It applies to fraud-operations analysts, automated fraud-review systems, and supporting investigation workflows.

The policy focuses on transaction characteristics that may indicate account compromise, unauthorized use, anomalous behavior, or elevated fraud risk. Policy conditions should be evaluated against authoritative transaction and customer data available to the investigation system.

2. Definitions

Card-Present (CP): A transaction in which the payment credential is physically presented or otherwise classified by the payment system as card present.

Card-Not-Present (CNP): A transaction in which the card is not physically presented, including qualifying e-commerce or remote transactions.

Anonymizing Proxy: A network endpoint identified by an approved intelligence source as a proxy, relay, Tor exit node, or other anonymization service.

Risk Level: The policy classification assigned to a transaction or investigation based on observed indicators and approved risk-scoring logic.

Secondary Fraud Review: A required review by a fraud analyst or designated fraud-review queue before the transaction or account is cleared.

3. General Fraud Risk Principles

Fraud investigations should use a combination of transaction facts, customer context, behavioral history, network indicators, and applicable policy requirements. No single indicator, by itself, necessarily establishes fraud.

Analysts must distinguish between an indicator, a policy trigger, and a final fraud determination. An indicator may increase concern; a policy trigger establishes a required control when all stated conditions are satisfied; a fraud determination may require additional evidence.

When evidence is incomplete or contradictory, the analyst should document the limitation and use the applicable review or escalation path rather than infer missing facts.

4. Transaction Velocity and Geographic Risk

4.1 Geographic Consistency - General Guidance

Review transaction geography when a customer exhibits rapid movement between materially separated locations. Consider the timestamps, transaction type, available location data, and whether the earlier transaction was card present.

The presence of geographically separated activity is an investigation signal. It becomes a policy trigger only when the conditions of a specific policy section are satisfied.

4.2 Impossible Travel Velocity

POLICY TRIGGER - SECTION 4.2
Any Card-Not-Present (CNP) transaction occurring within 2 hours of a physical card swipe in a different geographical continent must be treated as an Impossible Travel Velocity event and routed for elevated fraud review.

Required conditions: (1) the current transaction is classified as CNP; (2) a prior physical card swipe exists; (3) the elapsed time between the transactions is less than or equal to 2 hours; and (4) the locations are in different geographical continents.

Required handling: When all four conditions are satisfied, classify the event as a policy-triggered elevated-risk finding. If the approved risk score for the investigation exceeds the high-risk threshold defined in Section 7, apply the corresponding action requirements in Section 8.

Evidence required: Current transaction ID, current timestamp and location, current card-present classification, prior physical-card transaction timestamp and location, and enough geographic information to establish the continental difference.

4.3 Exceptions and Limitations

Do not classify a transaction under Section 4.2 when the prior transaction is not confirmed as physical card-present activity, when the time interval exceeds 2 hours, or when the available evidence does not establish different continents. Where geography cannot be determined reliably, route the case for manual review rather than infer a policy trigger.

5. Network and Anonymizing Proxy Risk

5.1 Network Risk - General Guidance

Network intelligence may identify IP addresses associated with known proxies, Tor exit nodes, or other anonymization services. Such information is an investigative signal and should be combined with transaction attributes and policy thresholds.

5.2 Proxy Intelligence Evidence

The preferred evidence for proxy-related review includes the originating IP address, the proxy classification returned by an approved intelligence source, transaction amount, timestamp, transaction type, and any relevant customer context available to the investigation.

5.3 Anonymizing Proxies

POLICY TRIGGER - SECTION 5.3
Transactions originating from known Tor exit nodes or proxy IPs carrying an amount higher than $500 trigger mandatory secondary fraud review.

Required conditions: (1) the originating network endpoint is classified as a known Tor exit node or proxy IP; and (2) the transaction amount is greater than $500.

Required handling: When both conditions are satisfied, route the transaction for mandatory secondary fraud review. The policy trigger does not, by itself, establish that the transaction is fraudulent.

Evidence required: Transaction ID, amount, originating IP address, network classification, and timestamp. Where available, retain the source or confidence indicator associated with the proxy classification.

5.4 Exceptions

A transaction below or equal to $500 does not meet the amount threshold in Section 5.3. If proxy status is unknown or unavailable, do not treat the transaction as a confirmed proxy event. Document the missing evidence.

6. Transaction Amount and Behavioral Anomalies

6.1 Customer Baseline

Where customer-level historical data is available, compare the current transaction amount against the customer average transaction amount and recent transaction behavior.

6.2 Unusual Amount Indicator

A transaction that is materially above the customer baseline should be recorded as a behavioral anomaly. This is an investigative finding unless a separate policy explicitly establishes a mandatory action or threshold.

Example calculation:

Amount ratio = Current transaction amount / Customer average transaction amount

Example: $1,450 / $120 = 12.08x the customer average transaction amount.

6.3 Interpretation

Amount anomalies should be considered together with other indicators such as geographic velocity, card-present status, proxy usage, merchant context, and recent transaction history. Avoid labeling a transaction as fraudulent solely because the amount is unusually high.

7. Risk Scoring and Classification

The investigation system may use an approved risk-scoring model or configured business scoring logic. The model or scoring mechanism is authoritative for any numerical risk score displayed to an analyst.

If the system provides a numerical score, display that score and the corresponding risk classification. If no numerical score is available, use a qualitative classification only when sufficient evidence exists.

Risk Level	Illustrative Score Range	General Handling
Low	0-39	Routine monitoring or closure when no policy trigger exists.
Medium	40-79	Review available evidence and apply any triggered controls.
High	80-100	Apply high-risk action requirements and escalate according to policy.

The ranges above are illustrative for this mock document. A production implementation should use the organization-approved risk model and configuration.

8. Required Actions and Escalation

8.1 High-Risk Transactions

HIGH-RISK ACTION RULE
When the approved risk assessment classifies an investigation as High Risk (risk level greater than 80), the required action is Immediate Card Freeze, subject to applicable operational controls and analyst confirmation.

8.2 Secondary Fraud Review

When Section 5.3 is triggered, the transaction must enter the secondary fraud review process. A secondary review may also be required by other policy sections or operational procedures.

8.3 Escalation

Escalate the investigation when the policy requires escalation, when a high-risk condition is established, or when evidence conflicts materially and cannot be resolved through available data sources.

8.4 Customer Notification

Customer notification may be initiated through approved channels when supported by the applicable operational procedure. The investigation system should expose a notification action only when the workflow permits it.

8.5 Action Traceability

Every displayed action should indicate the evidence or policy requirement that caused the action to be recommended. Do not expose an action as mandatory when the underlying policy condition is not satisfied.

9. Evidence and Analyst Documentation

Investigation records should preserve sufficient evidence to allow another analyst to understand why a finding or action was produced.

At minimum, record the Transaction ID, Customer ID, relevant transaction attributes, historical comparison evidence used, applicable policy section(s), triggered conditions, risk classification, and action taken or recommended.

Where an AI-assisted workflow is used, the analyst-facing result should distinguish source facts from interpretation and should provide policy section references for policy-driven conclusions.

10. Retrieval and Policy Interpretation Notes

10.1 Policy Retrieval

Semantic retrieval systems should retrieve the policy section most directly related to the observed transaction characteristic. Retrieval quality should be evaluated before a policy conclusion is generated.

10.2 Context Preservation

When a policy rule is retrieved, sufficient surrounding context should be retained to preserve definitions, thresholds, exceptions, and required actions. A single sentence fragment should not be used when surrounding language changes the meaning of the rule.

10.3 Corrective Retrieval

If the first retrieval result is irrelevant, incomplete, or missing a key threshold or condition, reformulate the retrieval query using concrete transaction evidence and retrieve again before reaching a conclusion.

10.4 Grounding

The investigation system must not invent a policy section, threshold, exception, or required action. If no directly applicable policy evidence is found, the system should state that the policy determination is inconclusive and identify the evidence gap.

10.5 Policy Citation Format

For analyst-facing output, cite the relevant section number and title, for example: Risk Manual §4.2 - Impossible Travel Velocity or Risk Manual §5.3 - Anonymizing Proxies.

Appendix A. Policy Evaluation Examples

Example A1 - Impossible Travel Velocity Trigger

Current transaction: CNP, London, 2026-07-29 10:15 UTC.
Prior transaction: physical card swipe, Minneapolis, 2026-07-29 09:45 UTC.
Time difference: 30 minutes.
Policy result: Section 4.2 conditions are satisfied if the locations are confirmed to be on different continents.

Example A2 - Anonymizing Proxy Trigger

Current transaction: $1,450.
Network: confirmed Tor exit node.

Policy result: Section 5.3 conditions are satisfied because the proxy condition is true and the transaction amount is greater than $500. Mandatory secondary fraud review is required.

Example A3 - Proxy But Below Threshold

Current transaction: $425.
Network: confirmed proxy IP.
Policy result: Section 5.3 is not triggered because the amount does not exceed $500. The proxy remains an investigative signal and may be relevant to other policy sections.

Example A4 - High-Risk Classification

Risk score: 88/100.
Policy result: The transaction is classified as High Risk under the illustrative scoring range in Section 7. Section 8.1 requires Immediate Card Freeze, subject to applicable operational controls and analyst confirmation.

Document Status: Synthetic demonstration policy. Created specifically for the U.S. Bank AI Hackathon fraud-investigation/RAG prototype. This document is not an official bank policy, does not represent legal or compliance advice, and must not be used for real operational decisions.